

Classification Problem

Inside this code, have

1. **Random Forest Classifier**
2. **KN-Neighbors Classifier**
3. **Support vector machine**
4. **Decision Tree**

Import Data and Libraries

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn import preprocessing
from sklearn.metrics import accuracy_score
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint
from IPython.display import Image
from sklearn.preprocessing import StandardScaler
```

```
In [3]: data=pd.read_csv("data.csv")
```

```
In [4]: data.head()
```

```
Out[4]:
```

| | baseline value | accelerations | fetal_movement | uterine_contractions | light_decelerations | severe_decelerations | prolongued_decelerations | abnormal_short_term |
|---|----------------|---------------|----------------|----------------------|---------------------|----------------------|--------------------------|---------------------|
| 0 | 120.0 | 0.000 | 0.0 | 0.000 | 0.000 | 0.0 | 0.0 | |
| 1 | 132.0 | 0.006 | 0.0 | 0.006 | 0.003 | 0.0 | 0.0 | |
| 2 | 133.0 | 0.003 | 0.0 | 0.008 | 0.003 | 0.0 | 0.0 | |
| 3 | 134.0 | 0.003 | 0.0 | 0.008 | 0.003 | 0.0 | 0.0 | |
| 4 | 132.0 | 0.007 | 0.0 | 0.008 | 0.000 | 0.0 | 0.0 | |

5 rows × 22 columns

```
In [5]: data.nunique()
```

```
Out[5]:
```

| | |
|--|-----|
| baseline value | 48 |
| accelerations | 20 |
| fetal_movement | 102 |
| uterine_contractions | 16 |
| light_decelerations | 16 |
| severe_decelerations | 2 |
| prolongued_decelerations | 6 |
| abnormal_short_term_variability | 75 |
| mean_value_of_short_term_variability | 57 |
| percentage_of_time_with_abnormal_long_term_variability | 87 |
| mean_value_of_long_term_variability | 249 |
| histogram_width | 154 |
| histogram_min | 109 |
| histogram_max | 86 |
| histogram_number_of_peaks | 18 |
| histogram_number_of_zeroes | 9 |
| histogram_mode | 88 |
| histogram_mean | 103 |
| histogram_median | 95 |
| histogram_variance | 133 |
| histogram_tendency | 3 |
| fetal_health | 3 |
| dtype: int64 | |

```
In [6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2126 entries, 0 to 2125
```

```
Data columns (total 22 columns):
```

| # | Column | Non-Null Count | Dtype |
|----|--|----------------|---------|
| 0 | baseline value | 2126 non-null | float64 |
| 1 | accelerations | 2126 non-null | float64 |
| 2 | fetal_movement | 2126 non-null | float64 |
| 3 | uterine_contractions | 2126 non-null | float64 |
| 4 | light_decelerations | 2126 non-null | float64 |
| 5 | severe_decelerations | 2126 non-null | float64 |
| 6 | prolongued_decelerations | 2126 non-null | float64 |
| 7 | abnormal_short_term_variability | 2126 non-null | float64 |
| 8 | mean_value_of_short_term_variability | 2126 non-null | float64 |
| 9 | percentage_of_time_with_abnormal_long_term_variability | 2126 non-null | float64 |
| 10 | mean_value_of_long_term_variability | 2126 non-null | float64 |
| 11 | histogram_width | 2126 non-null | float64 |
| 12 | histogram_min | 2126 non-null | float64 |
| 13 | histogram_max | 2126 non-null | float64 |
| 14 | histogram_number_of_peaks | 2126 non-null | float64 |
| 15 | histogram_number_of_zeroes | 2126 non-null | float64 |
| 16 | histogram_mode | 2126 non-null | float64 |
| 17 | histogram_mean | 2126 non-null | float64 |
| 18 | histogram_median | 2126 non-null | float64 |
| 19 | histogram_variance | 2126 non-null | float64 |
| 20 | histogram_tendency | 2126 non-null | float64 |
| 21 | fetal_health | 2126 non-null | float64 |

```
dtypes: float64(22)
```

```
memory usage: 365.5 KB
```

```
In [7]: data.isnull().sum()
```

```
Out[7]: baseline value      0
accelerations      0
fetal_movement     0
uterine_contractions 0
light_decelerations 0
severe_decelerations 0
prolongued_decelerations 0
abnormal_short_term_variability 0
mean_value_of_short_term_variability 0
percentage_of_time_with_abnormal_long_term_variability 0
mean_value_of_long_term_variability 0
histogram_width    0
histogram_min      0
histogram_max      0
histogram_number_of_peaks 0
histogram_number_of_zeroes 0
histogram_mode     0
histogram_mean     0
histogram_median   0
histogram_variance 0
histogram_tendency 0
fetal_health       0
dtype: int64
```

Exploratory Data Analysis

```
In [8]: data.columns
```

```
Out[8]: Index(['baseline value', 'accelerations', 'fetal_movement',
               'uterine_contractions', 'light_decelerations', 'severe_decelerations',
               'prolongued_decelerations', 'abnormal_short_term_variability',
               'mean_value_of_short_term_variability',
               'percentage_of_time_with_abnormal_long_term_variability',
               'mean_value_of_long_term_variability', 'histogram_width',
               'histogram_min', 'histogram_max', 'histogram_number_of_peaks',
               'histogram_number_of_zeroes', 'histogram_mode', 'histogram_mean',
               'histogram_median', 'histogram_variance', 'histogram_tendency',
               'fetal_health'],
              dtype='object')
```

```
In [9]: colours=["#f7b2b0", "#8f7198", "#003f5c"]
```

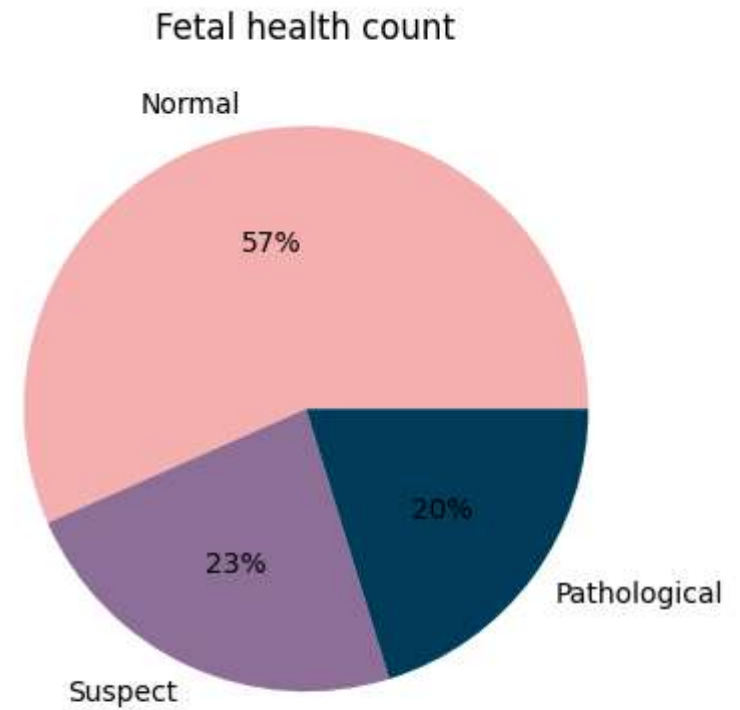

```
In [10]: total = data["fetal_health"].sum()
normal = total - 471
suspect = total - 1831
pathological = total - 1950

plt.figure(figsize = (14,8))
plt.subplot(121)
vis_fetal_health = data.fetal_health.value_counts().plot(figsize=(10, 5), kind="bar", color = colours)
plt.title("Fetal health count")
plt.xlabel("Fetal helth")
plt.ylabel("Cases")

plt.subplot(122)
plt.title("Fetal state")


vis_pie_fetal_health = plt.pie([normal, suspect, pathological], labels=["Normal", "Suspect", "Pathological"], colors = colours, a
plt.title("Fetal health count")

plt.show()
```



```
In [11]: plt.figure(figsize=(12,8))
sns.histplot(data=data, x='abnormal_short_term_variability', hue='fetal_health', palette=colours)
plt.suptitle('Distribution of Abnormal_hort_term_variability', size=16, y=1.01)
plt.tight_layout()
plt.show()
```

Distribution of Abnormal_hort_term_variability



Model Building

In [12]: `data.head()`

Out[12]:

| | baseline value | accelerations | fetal_movement | uterine_contractions | light_decelerations | severe_decelerations | prolongued_decelerations | abnormal_short_term |
|---|-------------------|---------------|----------------|----------------------|---------------------|----------------------|--------------------------|---------------------|
| 0 | 120.0 | 0.000 | 0.0 | 0.000 | 0.000 | 0.0 | 0.0 | |
| 1 | 132.0 | 0.006 | 0.0 | 0.006 | 0.003 | 0.0 | 0.0 | |
| 2 | 133.0 | 0.003 | 0.0 | 0.008 | 0.003 | 0.0 | 0.0 | |
| 3 | 134.0 | 0.003 | 0.0 | 0.008 | 0.003 | 0.0 | 0.0 | |
| 4 | 132.0 | 0.007 | 0.0 | 0.008 | 0.000 | 0.0 | 0.0 | |

5 rows × 22 columns

In [13]: *#assigning values to features as X and target as y*
`X=data.drop(["fetal_health"],axis=1)`
`y=data["fetal_health"]`

#Set up a standard scaler for the features
`col_names = list(X.columns)`
`s_scaler = preprocessing.StandardScaler()`
`X_df = s_scaler.fit_transform(X)`
`X_df = pd.DataFrame(X_df, columns=col_names)`

In [14]: `X`

Out[14]:

| | baseline value | accelerations | fetal_movement | uterine_contractions | light_decelerations | severe_decelerations | prolongued_decelerations | abnormal_short_t |
|------|-------------------|---------------|----------------|----------------------|---------------------|----------------------|--------------------------|------------------|
| 0 | 120.0 | 0.000 | 0.000 | 0.000 | 0.000 | 0.0 | 0.0 | |
| 1 | 132.0 | 0.006 | 0.000 | 0.006 | 0.003 | 0.0 | 0.0 | |
| 2 | 133.0 | 0.003 | 0.000 | 0.008 | 0.003 | 0.0 | 0.0 | |
| 3 | 134.0 | 0.003 | 0.000 | 0.008 | 0.003 | 0.0 | 0.0 | |
| 4 | 132.0 | 0.007 | 0.000 | 0.008 | 0.000 | 0.0 | 0.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 2121 | 140.0 | 0.000 | 0.000 | 0.007 | 0.000 | 0.0 | 0.0 | |
| 2122 | 140.0 | 0.001 | 0.000 | 0.007 | 0.000 | 0.0 | 0.0 | |
| 2123 | 140.0 | 0.001 | 0.000 | 0.007 | 0.000 | 0.0 | 0.0 | |
| 2124 | 140.0 | 0.001 | 0.000 | 0.006 | 0.000 | 0.0 | 0.0 | |
| 2125 | 142.0 | 0.002 | 0.002 | 0.008 | 0.000 | 0.0 | 0.0 | |

2126 rows × 21 columns

In [15]:

```
print(y)
```

```
0      2.0
1      1.0
2      1.0
3      1.0
4      1.0
...
2121    2.0
2122    2.0
2123    2.0
2124    2.0
2125    1.0
```

Name: fetal_health, Length: 2126, dtype: float64

In [16]:

```
X_train, X_test, y_train, y_test = train_test_split(X_df, y, test_size=0.25, random_state=42)
```

RandomForestClassifier

```
In [17]: from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn import svm

from sklearn.metrics import confusion_matrix
```

```
In [18]: model = RandomForestClassifier()
model.fit(X_train, y_train)
```

```
Out[18]: ▼ RandomForestClassifier
RandomForestClassifier()
```

```
In [19]: model_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(y_train, model_train_prediction)
print('Accuracy score of training data : ', training_data_accuracy)
```

Accuracy score of training data : 1.0

```
In [20]: model_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(y_test, model_test_prediction)
print('Accuracy score of test data : ', test_data_accuracy)
```


Accuracy score of test data : 0.9605263157894737

```
In [40]: comp = pd.DataFrame(
    {"Y_pred":model_test_prediction,
     "Y_test":y_test}
)

comp.to_csv('file1.csv')
```

```
In [22]: cm=confusion_matrix(model.predict(X_test),y_test)
sns.heatmap(cm,annot=True, fmt='g')
cm
```

```
Out[22]: array([[407, 10, 2],  
               [ 5, 70, 1],  
               [ 1, 2, 34]])
```



KN-Neighbors Classifier

```
In [23]: model_2=KNeighborsClassifier(n_neighbors=3)  
         model_2.fit(X_train,y_train)
```

```
Out[23]: ▼      KNeighborsClassifier  
         KNeighborsClassifier(n_neighbors=3)
```

```
In [24]: model_2_train_score=model_2.score(X_train,y_train)  
         print('Accuracy score of training data : ',model_2_train_score)
```

Accuracy score of training data : 0.9485570890840652


```
In [25]: modul_2_pred=model_2.predict(X_test)
```

```
In [26]: knn_pred_score=accuracy_score(y_test,modul_2_pred)
print('Accuracy score of test data : ',knn_pred_score)
```

Accuracy score of test data : 0.9229323308270677

```
In [27]: cm=confusion_matrix(model_2.predict(X_test),y_test)
sns.heatmap(cm,annot=True, fmt='g')
```

Out[27]: <Axes: >



```
In [27]:
```

Support vector machine

```
In [28]: model_3 = svm.SVC(kernel='rbf')
```

```
In [29]: # training the SVM model with training data
model_3.fit(X_train,y_train)
```

```
Out[29]: ▼ SVC
SVC()
```

```
In [30]: model_3_train_score=model_3.score(X_train,y_train)
print('Accuracy score of training data : ',model_3_train_score)
```

Accuracy score of training data : 0.9316185696361355

```
In [31]: modul_3_pred=model_3.predict(X_test)
```

```
In [32]: svm_pred_score=accuracy_score(y_test,modul_3_pred)
print('Accuracy score of test data : ',svm_pred_score)
```

Accuracy score of test data : 0.9172932330827067

Decision Tree

```
In [33]: model_4 = DecisionTreeClassifier()
```

```
In [34]: model_4.fit(X_train,y_train)
```

```
Out[34]: ▼ DecisionTreeClassifier
DecisionTreeClassifier()
```

```
In [35]: model_4_train_score=model_4.score(X_train,y_train)
print('Accuracy score of training data : ',model_4_train_score)
```

Accuracy score of training data : 1.0

```
In [36]: modul_4_pred=model_4.predict(X_test)
```

```
In [37]: dt_pred_score=accuracy_score(y_test,modul_4_pred)
print('Accuracy score of test data : ',dt_pred_score)
```

Accuracy score of test data : 0.9323308270676691

Results

```
In [38]: result=pd.DataFrame(columns=['Algorithm','Training-Result','Predictive-Result'])
```

```
In [39]: m=['KN-Neighbors Classifier', model_2_train_score, knn_pred_score]
result.loc[0]=m
m = ['Random Forest', training_data_accuracy, test_data_accuracy]
result.loc[1]=m
m = ['Support Vector Machine', model_3_train_score, svm_pred_score]
result.loc[2]=m
m = ['Decision Tree Classifier', model_4_train_score, dt_pred_score]
result.loc[3]=m

result
```

Out[39]:

| | Algorithm | Training-Result | Predictive-Result |
|--|-----------|-----------------|-------------------|
|--|-----------|-----------------|-------------------|

| | | | |
|---|--------------------------|----------|----------|
| 0 | KN-Neighbors Classifier | 0.948557 | 0.922932 |
| 1 | Random Forest | 1.000000 | 0.960526 |
| 2 | Support Vector Machine | 0.931619 | 0.917293 |
| 3 | Decision Tree Classifier | 1.000000 | 0.932331 |

In [39]: